

# Term Deposit Prediction

Here we are given a dataset of a bank in Portugal. Our job here is to use the existing data to gather insights and build a ML model that will determine how likely a person will subscribe to a term deposit. This will help us assess the marketing campaign that will be most effective. We will also perform EDA to answer questions along the way.

```
In [1]: #loading the dataset
```

```
import pandas as pd
```

```
data=pd.read_excel('C:\\Users\\sujoydutta\\Desktop\\Data analysis\\Python\\Projects\\term deposit.xlsx')
data.head()
```

```
Out[1]:
```

	age	job	marital	education	defaulter?	yearly_balance	housing_loan	personal_loan	contacted_via
0	58	management	married	tertiary	no	2143	yes	no	unknown
1	44	technician	single	secondary	no	29	yes	no	unknown
2	33	entrepreneur	married	secondary	no	2	yes	yes	unknown
3	47	blue-collar	married	unknown	no	1506	yes	no	unknown
4	33	unknown	single	unknown	no	1	no	no	unknown

```
In [2]: #examining the dataset
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 45211 entries, 0 to 45210
```

```
Data columns (total 17 columns):
```

#	Column	Non-Null Count	Dtype
0	age	45211 non-null	int64
1	job	45211 non-null	object
2	marital	45211 non-null	object
3	education	45211 non-null	object
4	defaulter?	45211 non-null	object
5	yearly_balance	45211 non-null	int64
6	housing_loan	45211 non-null	object
7	personal_loan	45211 non-null	object
8	contacted_via	45211 non-null	object
9	day	45211 non-null	int64
10	month	45211 non-null	object
11	duration_call_seconds	45211 non-null	int64
12	num_times_contacted	45211 non-null	int64
13	days_bt看_contact	45211 non-null	int64
14	times_contact_before	45211 non-null	int64
15	pc_outcome	45211 non-null	object
16	purchased?	45211 non-null	object

```
dtypes: int64(7), object(10)
```

```
memory usage: 5.9+ MB
```

```
In [3]: import matplotlib.pyplot as plt
```

```
import seaborn as sns
```


```
# Pie chart to see the distribution of outcome
```

```
outcome_counts = data['purchased?'].value_counts()
```

```
plt.figure(figsize=(8, 8))
```

```
plt.pie(outcome_counts, labels=outcome_counts.index, autopct='%1.1f%%', startangle=140)
plt.title('What percentage of people bought term deposit?')
plt.show()
```

What percentage of people bought term deposit?




```
In [4]: # Which month has highest conversion and which month has lowest conversion

df_yes = data[data['purchased?'] == 'yes']

plt.figure(figsize=(12, 6))
sns.countplot(data=df_yes, x='month', order=df_yes['month'].value_counts().index)
plt.title('Month-wise Distribution of Outcome "Yes"')
plt.xlabel('Month')
plt.ylabel('Count')
plt.show()
```

Month-wise Distribution of Outcome "Yes"




```
In [5]: # what percentage of loan customers are previous defaulters?

df_yes = data[data['purchased?'] == 'yes']

default_counts = df_yes['defaulter?'].value_counts()

plt.figure(figsize=(8, 8))
plt.pie(default_counts, labels=default_counts.index, autopct='%1.1f%%', startangle=140,
plt.title('Percentage of loan customers previous defaulters')
plt.show()
```

## Percentage of loan customers previous defaulters




```
In [6]: # what percentage of loan customers have Housing loans?

df_yes = data[data['purchased?'] == 'yes']

default_counts = df_yes['housing_loan'].value_counts()

plt.figure(figsize=(8, 8))
plt.pie(default_counts, labels=default_counts.index, autopct='%1.1f%%', startangle=140,
plt.title('Percentage of loan customers who have housing loans')
plt.show()
```

## Percentage of loan customers who have housing loans




```
In [7]: # what percentage of loan customers have personal loans?

df_yes = data[data['purchased?'] == 'yes']

default_counts = df_yes['personal_loan'].value_counts()

plt.figure(figsize=(8, 8))
plt.pie(default_counts, labels=default_counts.index, autopct='%1.1f%%', startangle=140,
plt.title('Percentage of loan customers who have personal loans')
plt.show()
```

## Percentage of loan customers who have personal loans




```
In [8]: # Share of each marital status that purchased the term deposit

df_yes = data[data['purchased?'] == 'yes']

default_counts = df_yes['marital'].value_counts()

plt.figure(figsize=(8, 8))
plt.pie(default_counts, labels=default_counts.index, autopct='%1.1f%%', startangle=140,
plt.title('Percentage of loan customers by marital status')
plt.show()
```

Percentage of loan customers by marital status




```
In [9]: # Share of each education level that purchased the term deposit

df_yes = data[data['purchased?'] == 'yes']

default_counts = df_yes['education'].value_counts()

plt.figure(figsize=(8, 8))
plt.pie(default_counts, labels=default_counts.index, autopct='%1.1f%%', startangle=140,
plt.title('Percentage of loan customers by education level')
plt.show()
```

Percentage of loan customers by education level



```
In [10]: # Share of each profession type that purchased the term deposit


df_yes = data[data['purchased?'] == 'yes']

default_counts = df_yes['job'].value_counts()

plt.figure(figsize=(8, 8))
plt.pie(default_counts, labels=default_counts.index, autopct='%1.1f%%', startangle=140,
plt.title('Percentage of loan customers by job type')
plt.show()
```

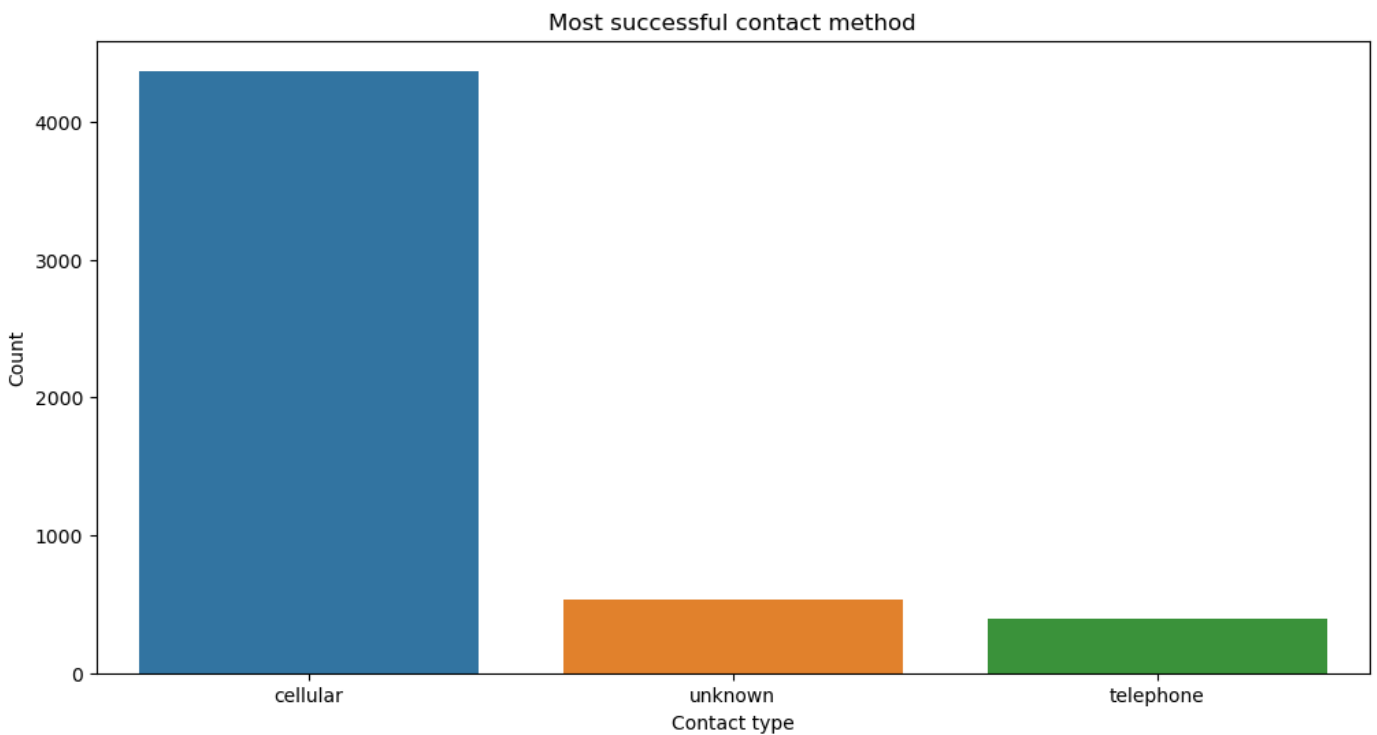


Percentage of loan customers by job type



```
In [11]: # Which medium of communication has highest conversion and which one has lowest conversion
df_yes = data[data['purchased?'] == 'yes']

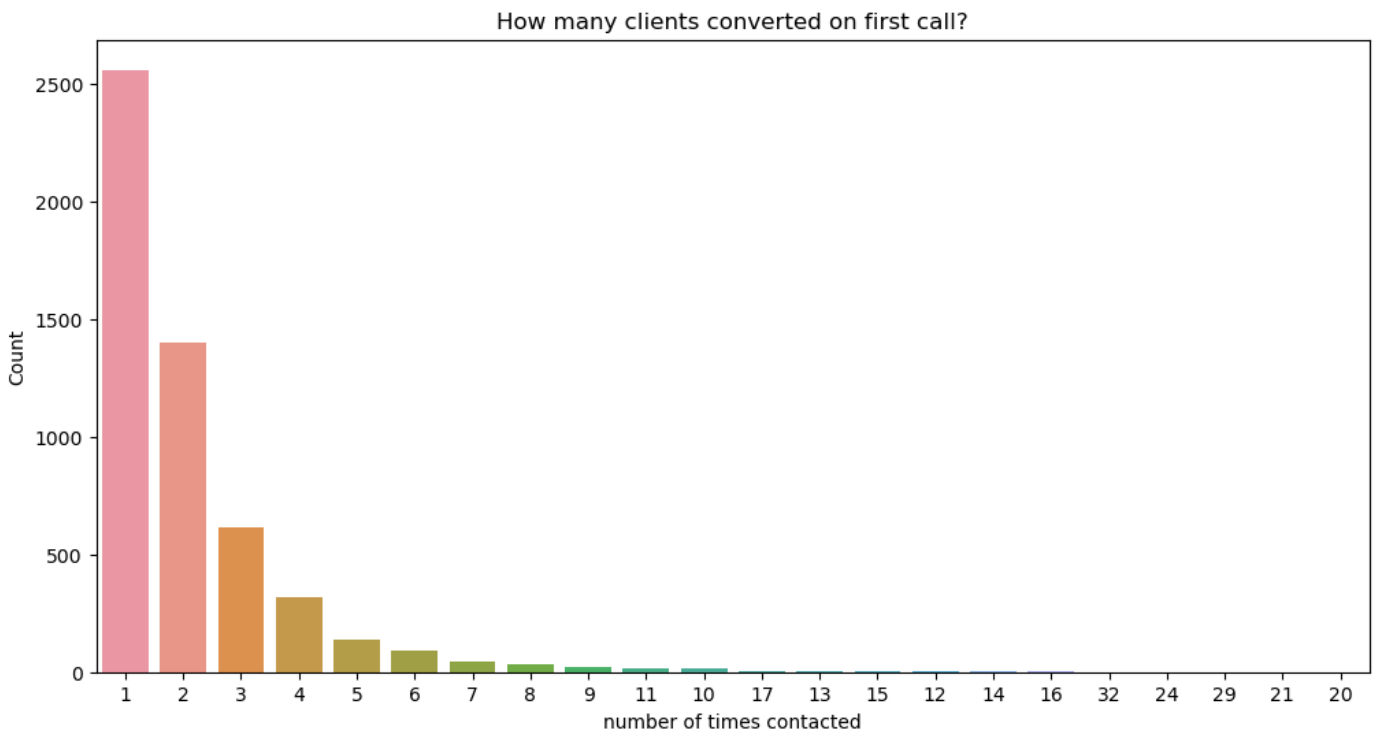
plt.figure(figsize=(12, 6))
sns.countplot(data=df_yes, x='contacted_via', order=df_yes['contacted_via'].value_counts)
plt.title('Most successful contact method')
plt.xlabel('Contact type')
plt.ylabel('Count')
plt.show()
```



```
In [12]: # How many clients converted on the first call?

df_yes = data[data['purchased?'] == 'yes']

plt.figure(figsize=(12, 6))
sns.countplot(data=df_yes, x='num_times_contacted', order=df_yes['num_times_contacted'].
plt.title('How many clients converted on first call?')
plt.xlabel('number of times contacted')
plt.ylabel('Count')
plt.show()
```



```
In [13]: # Seeing the correlation among numerical variables
numerical_columns = data.select_dtypes(include=['int64', 'float64']).columns
numerical_data = data[numerical_columns]
```

```
corr_matrix = numerical_data.corr()
```

```
plt.figure(figsize=(12, 8))
```

```
sns.heatmap(corr_matrix, annot=True, fmt='.2f', cmap='coolwarm', vmin=-1, vmax=1, linewidths=0.5)
```

```
plt.title('Correlation Matrix of Numerical Variables')
```

```
plt.show()
```



In [14]: *# What should be the outcome of previous contact to get a term deposit conversion>*

```
df_yes = data[data['purchased?'] == 'yes']
```

```
plt.figure(figsize=(12, 6))
```


```
sns.countplot(data=df_yes, x='pc_outcome', order=df_yes['pc_outcome'].value_counts().index)
```

```
plt.title('Most successful contact method')
```

```
plt.xlabel('Previous contact outcome')
```

```
plt.ylabel('Count')
```

```
plt.show()
```




```
In [15]: # Viewing age distribution of the campaign target
plt.subplot(1, 2, 2)
sns.kdeplot(data['age'], shade=True, color='skyblue')
plt.title('Age Distribution - KDE')
plt.xlabel('Age')
plt.ylabel('Density')

plt.tight_layout()
plt.show()
```


C:\Users\sujoydutta\AppData\Local\Temp\ipykernel\_12728\1407996061.py:3: FutureWarning:  
`shade` is now deprecated in favor of `fill`; setting `fill=True`.  
This will become an error in seaborn v0.14.0; please update your code.

```
sns.kdeplot(data['age'], shade=True, color='skyblue')
```



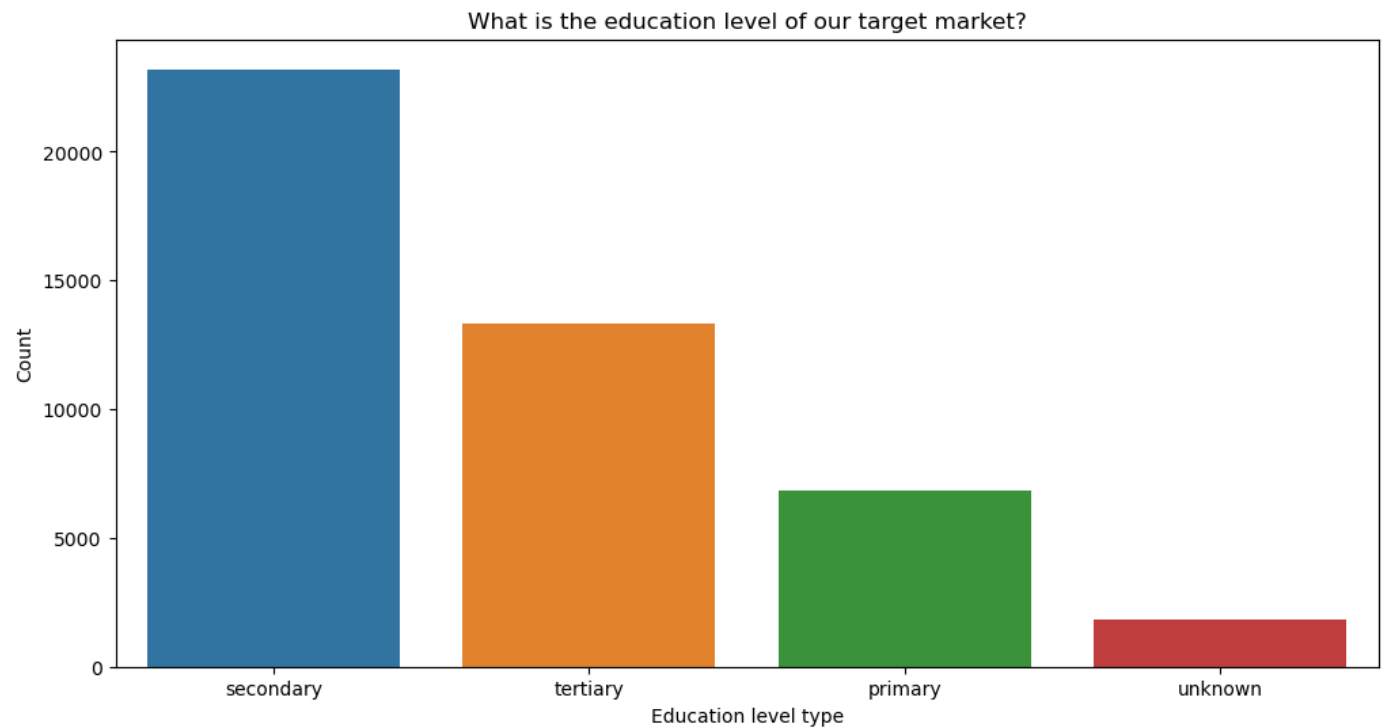
```
In [16]: # Marital status distribution of the target group

plt.figure(figsize=(12, 6))
sns.countplot(data=data, x='marital', order=data['marital'].value_counts().index)
plt.title('What is the marital status of our target market?')
plt.xlabel('Marital status type')
plt.ylabel('Count')
plt.show()
```



In [17]: *# level of education distribution of the target group*


```
plt.figure(figsize=(12, 6))
sns.countplot(data=data, x='education', order=data['education'].value_counts().index)
plt.title('What is the education level of our target market?')
plt.xlabel('Education level type')
plt.ylabel('Count')
plt.show()
```



In [18]: *# Type of profession the target group engages in*

```
plt.figure(figsize=(15, 8))
sns.countplot(data=data, x='job', order=data['job'].value_counts().index)
plt.title('What is the job type of our target market?')
plt.xlabel('Profession type')
plt.ylabel('Count')
plt.show()
```

What is the job type of our target market?



```
In [19]: # Filtering the dataset based on the outcome
purchased = data[data['purchased?'] == 'yes']
not_purchased = data[data['purchased?'] == 'no']
```

```
In [20]: import numpy as np
from scipy import stats

# Function to perform z-test
def z_test(group1, group2, variable):
    mean1 = group1[variable].mean()
    mean2 = group2[variable].mean()
    std1 = group1[variable].std()
    std2 = group2[variable].std()
    n1 = len(group1[variable])
    n2 = len(group2[variable])

    z = (mean1 - mean2) / np.sqrt((std1**2 / n1) + (std2**2 / n2))
    p = 2 * (1 - stats.norm.cdf(abs(z)))
    return z, p
```

```
In [21]: # Performing z-test for each variable
variables = ['num_times_contacted', 'days_btw_contact', 'yearly_balance', 'duration_call_s

results = {}
for var in variables:
    z, p = z_test(purchased, not_purchased, var)
    results[var] = {'z-score': z, 'p-value': p}

results_df = pd.DataFrame(results).T
print(results_df)
```

	z-score	p-value
num_times_contacted	-22.800741	0.000000
days_btw_contact	18.943484	0.000000
yearly_balance	9.933545	0.000000
duration_call_seconds	57.514127	0.000000
times_contact_before	18.117970	0.000000
age	4.318318	0.000016

```
In [22]: #visualizing the difference
variables = ['duration_call_seconds', 'num_times_contacted', 'days_bt看_contact', 'yearly_
means_purchased = purchased[variables].mean()
means_not_purchased = not_purchased[variables].mean()

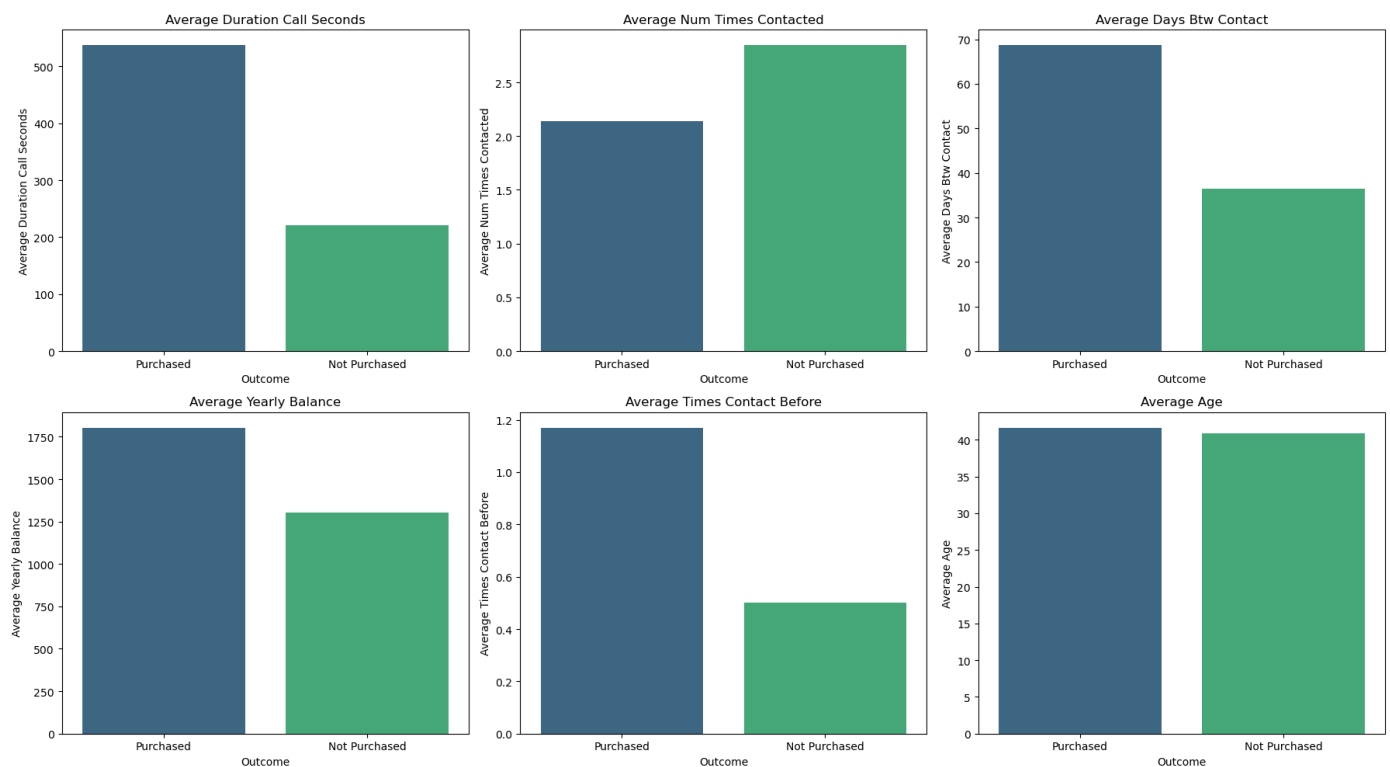
means_df = pd.DataFrame({
    'Purchased': means_purchased,
    'Not Purchased': means_not_purchased
}).T

plt.figure(figsize=(18, 10))

for i, var in enumerate(variables):
    plt.subplot(2, 3, i + 1)
    sns.barplot(data=means_df[[var]].reset_index(), x='index', y=var, palette='viridis')
    plt.title(f'Average {var.replace("_", " ").title()}')
    plt.xlabel('Outcome')
    plt.ylabel(f'Average {var.replace("_", " ").title()}')

if len(variables) < 6:
    plt.subplot(2, 3, 6).axis('off')

plt.tight_layout()
plt.show()
```



```
In [23]: # Function to identify outliers using Z-score
def identify_outliers_zscore(df, threshold=3):
    outliers = (np.abs(stats.zscore(df)) > threshold).any(axis=1)
    return outliers
```

```
In [24]: # Identifying outliers in the numerical columns
numerical_columns = ['age', 'yearly_balance', 'duration_call_seconds', 'num_times_contacte

outliers = identify_outliers_zscore(data[numerical_columns])
outliers
```



```
Out[24]: 0      False
1      False
2      False
3      False
4      False
...
45206   False
45207   False
45208    True
45209   False
45210    True
Length: 45211, dtype: bool
```

```
In [25]: # Analyzing the effect of outliers
data_no_outliers = data[~outliers]
data_with_outliers = data[outliers]
```

```
In [26]: # Comparing the datasets
print(f'Original dataset size: {data.shape[0]}')
print(f'Dataset size without outliers: {data_no_outliers.shape[0]}')
print(f'Dataset size with outliers: {data_with_outliers.shape[0]}')
```


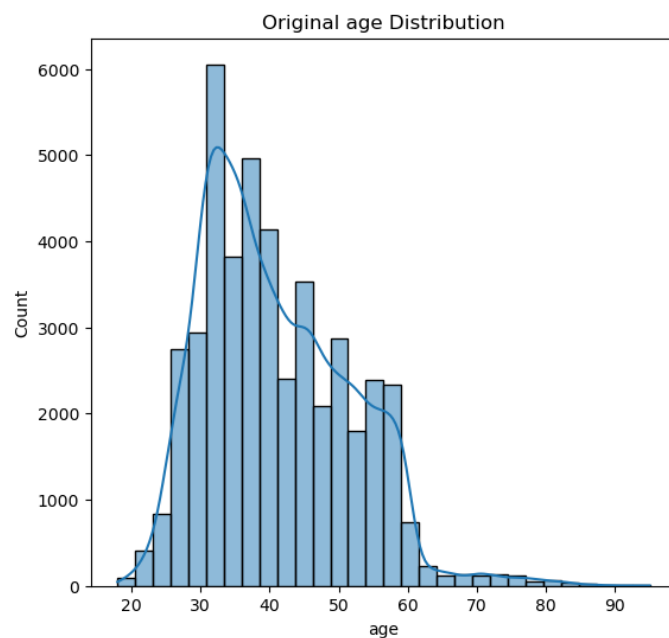
```
Original dataset size: 45211
Dataset size without outliers: 40209
Dataset size with outliers: 5002
```

```
In [27]: # Visualizing the distributions before and after outlier removal
for col in numerical_columns:
    plt.figure(figsize=(14, 6))


    plt.subplot(1, 2, 1)
    sns.histplot(data[col], bins=30, kde=True)
    plt.title(f'Original {col} Distribution')

    plt.subplot(1, 2, 2)
    sns.histplot(data_no_outliers[col], bins=30, kde=True)
    plt.title(f'{col} Distribution without Outliers')


    plt.show()
```




Original yearly\_balance Distribution




yearly\_balance Distribution without Outliers




Original duration\_call\_seconds Distribution




duration\_call\_seconds Distribution without Outliers

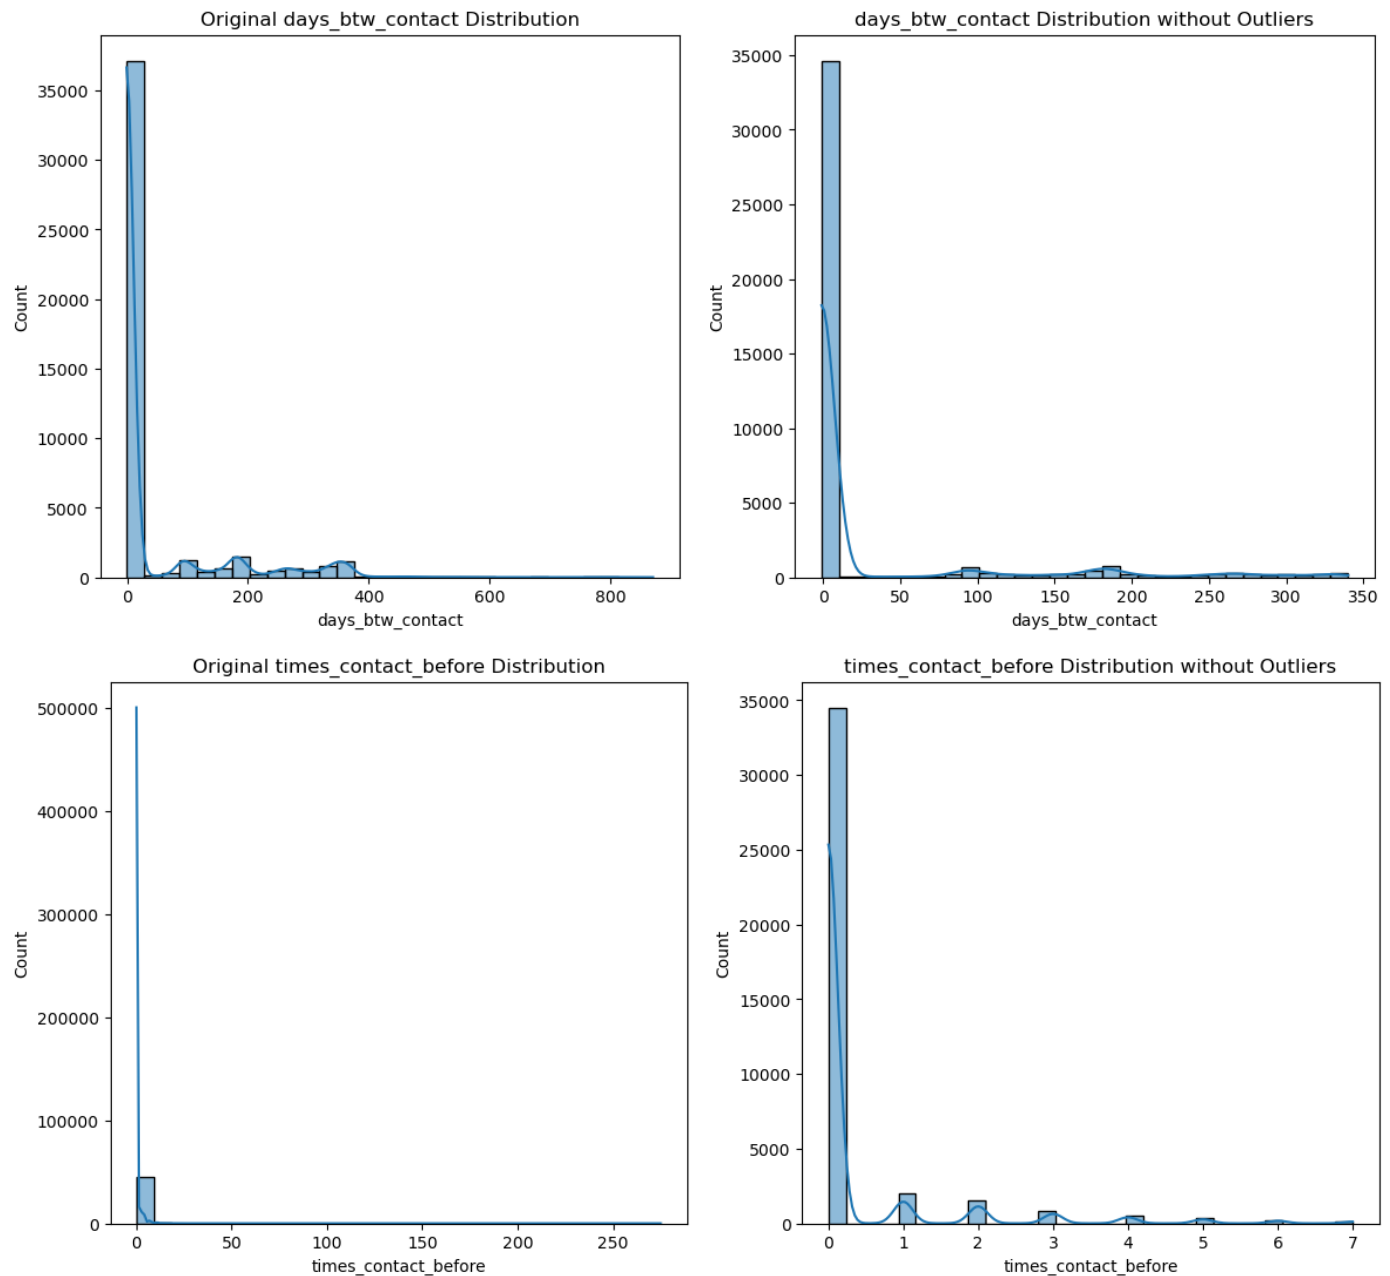


Original num\_times\_contacted Distribution



num\_times\_contacted Distribution without Outliers





```
In [28]: # Replacing outliers with median values
for col in numerical_columns:
    median = data[col].median()
    data.loc[outliers, col] = median
```

```
In [29]: # Checking if outliers are replaced
print(data.describe())
```

	age	yearly_balance	day	duration_call_seconds \
count	45211.000000	45211.000000	45211.000000	45211.000000
mean	40.476101	1013.160249	15.806419	229.621773
std	9.526543	1636.907495	8.322476	180.417341
min	18.000000	-6847.000000	1.000000	0.000000
25%	33.000000	99.000000	8.000000	112.000000
50%	39.000000	448.000000	16.000000	180.000000
75%	47.000000	1161.000000	21.000000	283.000000
max	72.000000	10483.000000	31.000000	1030.000000

	num_times_contacted	days_bt看_contact	times_contact_before
count	45211.000000	45211.000000	45211.000000
mean	2.412953	22.287364	0.313618
std	1.846032	68.173144	1.002829
min	1.000000	-1.000000	0.000000
25%	1.000000	-1.000000	0.000000

50%	2.000000	-1.000000	0.000000
75%	3.000000	-1.000000	0.000000
max	12.000000	340.000000	7.000000

```
In [30]: from statsmodels.stats.outliers_influence import variance_inflation_factor
# Calculating VIF for each feature
X = data[numerical_columns]
vif_data = pd.DataFrame()
vif_data['Feature'] = X.columns
vif_data['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]

print(vif_data)
```

	Feature	VIF
0	age	4.309325
1	yearly_balance	1.405196
2	duration_call_seconds	2.403441
3	num_times_contacted	2.549530
4	days_btw_contact	2.322813
5	times_contact_before	2.309668

```
In [31]: # Define categorical and numerical columns
categories = ['job', 'marital', 'education', 'day', 'defaulter?', 'housing_loan', 'person'
numerical_columns = ['duration_call_seconds', 'num_times_contacted', 'days_btw_contact',
```

```
In [32]: # Define target column
target = 'purchased?'

# Preparing feature matrix and target vector
X = data[categories + numerical_columns]
y = data[target].map({'yes': 1, 'no': 0})
```

```
In [35]: from sklearn.preprocessing import StandardScaler
import category_encoders as ce
# Defining preprocessing steps
numeric_transformer = StandardScaler()
categorical_transformer = ce.BinaryEncoder(cols=categories)
```

```
In [36]: from sklearn.compose import ColumnTransformer

# Combining preprocessing steps
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numerical_columns),
        ('cat', categorical_transformer, categories)
    ])

```

```
In [37]: from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
# Defining the model pipeline
model = Pipeline(steps=[('preprocessor', preprocessor),
                        ('classifier', LogisticRegression())])
```

```
In [38]: from sklearn.model_selection import train_test_split

# Splitting the data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=4
```

```
In [39]: # Fitting the model
model.fit(X_train, y_train)
```

C:\Users\sujoydutta\anaconda\Lib\site-packages\sklearn\linear\_model\\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

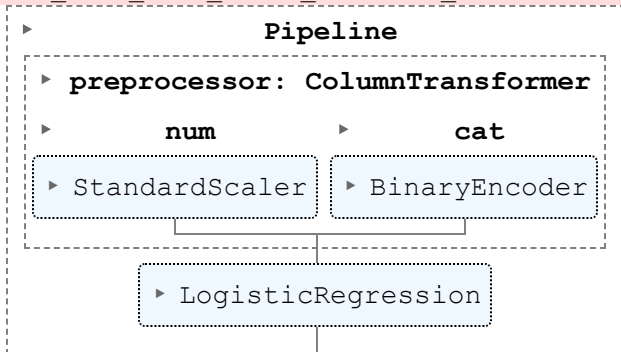
<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Out[39]:



In [40]:

```
# Predicting on the test set
y_pred = model.predict(X_test)
```

In [41]:

```
# Evaluating the model
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

print(f'Accuracy: {accuracy_score(y_test, y_pred)}')
print('Classification Report:')
print(classification_report(y_test, y_pred))
print('Confusion Matrix:')
print(confusion_matrix(y_test, y_pred))
```

Accuracy: 0.8948066884897815

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.98	0.94	9950
1	0.64	0.28	0.39	1353
accuracy			0.89	11303
macro avg	0.77	0.63	0.67	11303
weighted avg	0.88	0.89	0.88	11303

Confusion Matrix:

```
[[9736 214]
 [ 975 378]]
```

In [42]:

```
# Example with Logistic Regression
model = Pipeline(steps=[('preprocessor', preprocessor),
                        ('classifier', LogisticRegression(class_weight='balanced'))])
```

In [43]:

```
# Fitting the model
model.fit(X_train, y_train)
```

C:\Users\sujoydutta\anaconda\Lib\site-packages\sklearn\linear\_model\\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

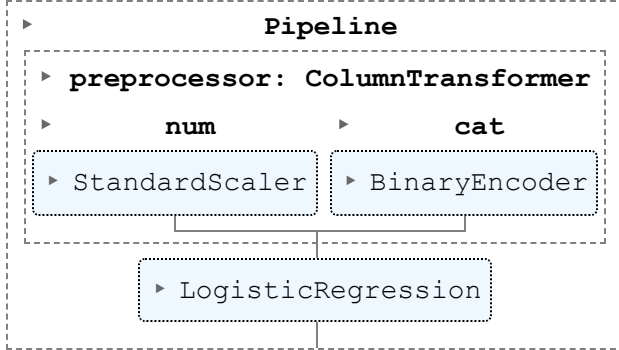
<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Out[43]:



```
In [44]: # Predicting on the test set
y_pred = model.predict(X_test)
```

```
In [45]: # Evaluating the model

print(f'Accuracy: {accuracy_score(y_test, y_pred)}')
print('Classification Report:')
print(classification_report(y_test, y_pred))
print('Confusion Matrix:')
print(confusion_matrix(y_test, y_pred))
```

```
Accuracy: 0.7860744934973016
Classification Report:
              precision    recall  f1-score   support


     0       0.96      0.79      0.87      9950
     1       0.33      0.75      0.46      1353

 accuracy          0.79      11303
 macro avg       0.64      0.77      0.66      11303
 weighted avg    0.88      0.79      0.82      11303

Confusion Matrix:
[[7868 2082]
 [ 336 1017]]
```

```
In [46]: from sklearn.ensemble import RandomForestClassifier
# Example with Random Forest Classifier
model = Pipeline(steps=[('preprocessor', preprocessor),
                        ('classifier', RandomForestClassifier(class_weight='balanced', r
```

```
In [47]: # Fitting the model
model.fit(X_train, y_train)
```



```
In [48]: # Evaluating the model

print(f'Accuracy: {accuracy_score(y_test, y_pred)}')
print('Classification Report:')
print(classification_report(y_test, y_pred))
```

```
print('Confusion Matrix:')
print(confusion_matrix(y_test, y_pred))
```

Accuracy: 0.7860744934973016

Classification Report:

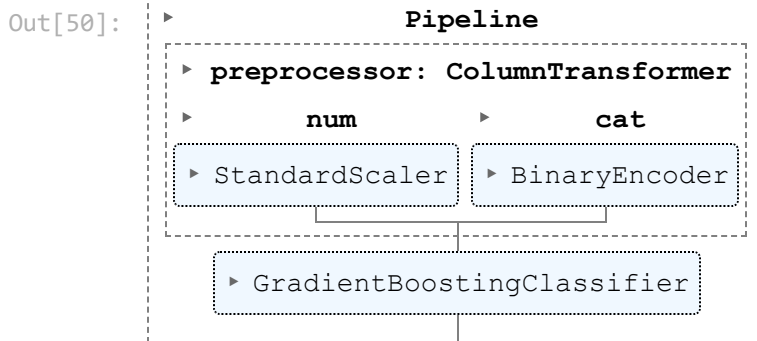
	precision	recall	f1-score	support
0	0.96	0.79	0.87	9950
1	0.33	0.75	0.46	1353
accuracy			0.79	11303
macro avg	0.64	0.77	0.66	11303
weighted avg	0.88	0.79	0.82	11303

Confusion Matrix:

```
[[7868 2082]
 [ 336 1017]]
```

```
In [49]: from sklearn.ensemble import GradientBoostingClassifier
# Example with Gradient Boosting Classifier
model = Pipeline(steps=[('preprocessor', preprocessor),
                        ('classifier', GradientBoostingClassifier(random_state=42))])
```

```
In [50]: # Fitting the model
model.fit(X_train, y_train)
```



```
In [51]: # Predicting on the test set
y_pred = model.predict(X_test)
```

```
In [52]: # Evaluating the model

print(f'Accuracy: {accuracy_score(y_test, y_pred)}')
print('Classification Report:')
print(classification_report(y_test, y_pred))
print('Confusion Matrix:')
print(confusion_matrix(y_test, y_pred))
```

Accuracy: 0.8956029372732903

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.98	0.94	9950
1	0.65	0.27	0.38	1353
accuracy			0.90	11303
macro avg	0.78	0.63	0.66	11303
weighted avg	0.88	0.90	0.88	11303

Confusion Matrix:

```
[[9757 193]
 [ 987 366]]
```

```
In [86]: # Defining the parameter grid for GridSearchCV
```

```

param_grid = {
    'classifier__n_estimators': [50, 100, 200],
    'classifier__learning_rate': [0.1, 0.05, 0.01],
    'classifier__max_depth': [3, 4, 5],
    'classifier__min_samples_split': [2, 5, 10],
    'classifier__min_samples_leaf': [1, 2, 4],
    'classifier__subsample': [0.6, 0.7, 0.8],
    'classifier__max_features': [None, 'sqrt', 'log2']
}

```

```

In [87]: #building the pipeline
pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                             ('classifier', GradientBoostingClassifier(random_state=42))])

```

```

In [88]: from sklearn.experimental import enable_halving_search_cv
from sklearn.model_selection import HalvingGridSearchCV

# Performing GridSearchCV to find the best parameters
grid_search = HalvingGridSearchCV(pipeline, param_grid, cv=5, verbose=1, n_jobs=-1)
grid_search.fit(X, y)

```

```

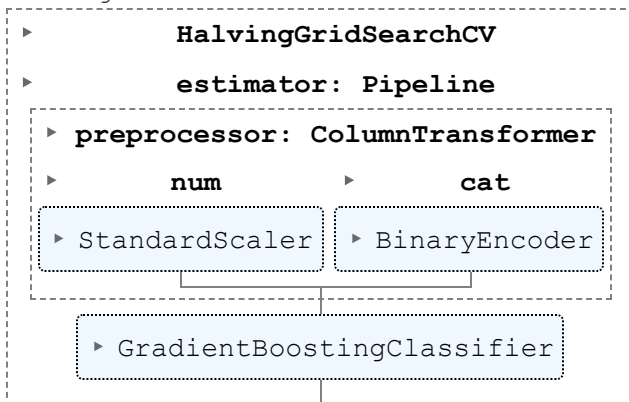
n_iterations: 8
n_required_iterations: 8
n_possible_iterations: 8
min_resources_: 20
max_resources_: 45211
aggressive_elimination: False
factor: 3
-----
iter: 0
n_candidates: 2187
n_resources: 20
Fitting 5 folds for each of 2187 candidates, totalling 10935 fits
-----
iter: 1
n_candidates: 729
n_resources: 60
Fitting 5 folds for each of 729 candidates, totalling 3645 fits
-----
iter: 2
n_candidates: 243
n_resources: 180
Fitting 5 folds for each of 243 candidates, totalling 1215 fits
-----
iter: 3
n_candidates: 81
n_resources: 540
Fitting 5 folds for each of 81 candidates, totalling 405 fits
-----
iter: 4
n_candidates: 27
n_resources: 1620
Fitting 5 folds for each of 27 candidates, totalling 135 fits
-----
iter: 5
n_candidates: 9
n_resources: 4860
Fitting 5 folds for each of 9 candidates, totalling 45 fits
-----
iter: 6
n_candidates: 3
n_resources: 14580
Fitting 5 folds for each of 3 candidates, totalling 15 fits
-----
iter: 7

```



```
n_candidates: 1
n_resources: 43740
Fitting 5 folds for each of 1 candidates, totalling 5 fits
```

Out[88]:



In [89]:

```
# Printing the best parameters
print("Best parameters found:")
print(grid_search.best_params_)
```

Best parameters found:

```
{'classifier__learning_rate': 0.01, 'classifier__max_depth': 5, 'classifier__max_features': 'log2', 'classifier__min_samples_leaf': 2, 'classifier__min_samples_split': 5, 'classifier__n_estimators': 50, 'classifier__subsample': 0.6}
```

In [90]:

```
# Printing the best score
print("Best CV score:")
print(grid_search.best_score_)
```

Best CV score:

```
0.883254061265711
```

In [92]:

```
# Best parameters from GridSearchCV
best_params = {
    'learning_rate': 0.01,
    'max_depth': 5,
    'max_features': 'log2',
    'min_samples_leaf': 2,
    'min_samples_split': 5,
    'n_estimators': 50,
    'subsample': 0.6
}
```

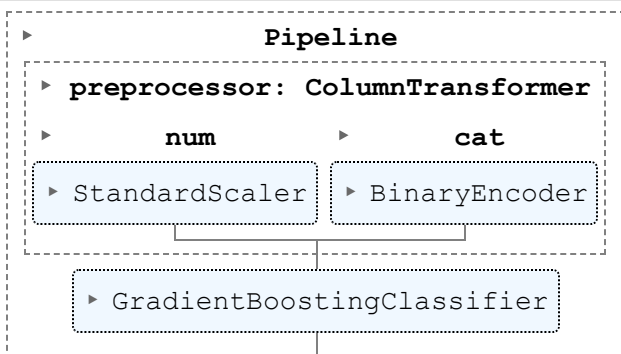
In [93]:

```
# Created the final model with the best parameters
final_model = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', GradientBoostingClassifier(**best_params, random_state=42))
])
```

In [94]:

```
# Fitting the model
final_model.fit(X_train, y_train)
```

Out[94]:



```
In [95]: # Printing the training accuracy
train_score = final_model.score(X_train, y_train)
print("Training Accuracy: ", train_score)
```

Training Accuracy: 0.8839211985372184

```
In [96]: # Printing the test accuracy
test_score = final_model.score(X_test, y_test)
print("Test Accuracy: ", test_score)
```

Test Accuracy: 0.8802972662125099

```
In [97]: # Evaluating the model

print(f'Accuracy: {accuracy_score(y_test, y_pred)}')
print('Classification Report:')
print(classification_report(y_test, y_pred))
print('Confusion Matrix:')
print(confusion_matrix(y_test, y_pred))
```

Accuracy: 0.8956029372732903

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.98	0.94	9950
1	0.65	0.27	0.38	1353
accuracy			0.90	11303
macro avg	0.78	0.63	0.66	11303
weighted avg	0.88	0.90	0.88	11303

Confusion Matrix:

```
[[9757 193]
 [ 987 366]]
```


```
In [98]: #getting the importance of variables
feature_importances = final_model.named_steps['classifier'].feature_importances_
```

```
In [99]: # Lets combine the names of numerical and categorical features
feature_names = (
    final_model.named_steps['preprocessor']
    .transformers_[0][2] + # Numerical features
    final_model.named_steps['preprocessor']
    .transformers_[1][1] # OneHotEncoder feature names
    .get_feature_names_out(final_model.named_steps['preprocessor']
    .transformers_[1][2]).tolist()
)
```

```
In [100]: # Creating a DataFrame for better visualization
feature_importance_df = pd.DataFrame({
    'Feature': feature_names,
    'Importance': feature_importances
}).sort_values(by='Importance', ascending=False)
```

```
In [101]: # Plotting the feature importances
plt.figure(figsize=(10, 8))
plt.barh(feature_importance_df['Feature'], feature_importance_df['Importance'])
plt.xlabel('Feature Importance')
plt.ylabel('Feature')
plt.title('Feature Importance from Gradient Boosting Classifier')
plt.gca().invert_yaxis()
plt.show()
```

Feature Importance from Gradient Boosting Classifier



In [ ]: